Enhancing the causal predictive power in recurrent network models of neural dynamics

COSYNE 2025 • 3-046

Jiayi Zhang[1], Tatiana A. Engel[2]

[1] Lewis-Sigler Institute of Integrative Genomics, Princeton University, Princeton, NJ, USA
[2] Princeton Neuroscience Institute, Princeton University, Princeton, NJ, USA

Lewis-Sigler Institute for Integrative Genomics

Princeton Neuroscience Institute

NIH National Institutes of Health The BRAIN Initiative

Most updated poster version

Motivations

Optogenetics enables unprecedented precisions for ***observing*** and ***perturbing neural circuits in actions***.
Designing perturbations that test current theories or enhance insights into circuit mechanisms ***remains an open challenge***[1].
RNNs that replicate observed neural dynamics might serve as ***"digital twins"*** of brain circuits[2].
Can these digital twins help with designing effective perturbation patterns for causal validation of circuit mechnaisms?

Twins cannot predict random perturbations

Dream:
large effect size; accurate predictions

Reality:
large effect size;
inaccurate predictions

Ground-truth (GT): 200-node RNN, trained to perform delayed decision-making task in Daie et al., 2021[1].
Twins (500 replicates): 200-node RNNs, trained on the same task as GT and required to replicate GT neural dynamics.
Paired (random) perturbations: inject the same (out-of-distribution) stimulus to twins and GT; measure R^2 between twin and GT trajectories.

Likely cause: Dim. of data << Dim. of observed units. ***Low-D data cannot sufficiently constrain the full connectivity*** of twin models.

Fitting models of varying dimensionalty

Data uniquely constrain RNNs in low-D latent space

x-axis: R^2(GT before & after perturbation)
y-axis: R^2(GT vs. Twin after perturbation)

(Top) Twins can predict GT trajectories with high accuracies when perturbations are directed along the latent directions!
(Left) All twins harbor the same low-D latent structure, despite their discrepancies with the GT in the full space.

Latent circuit inference

Latent circuits (100 replicates, size 4-10 nodes): small RNN, trained on the same task as GT and generates neural dynamics x. An orthonormal embedding matrix Q is learnt simultaneously Qx approximates y.

Under successful inference, low-D connectivity is a latent structure within the full connectivity:

$$Q^T W_{rec} Q \sim w_{rec}$$

Fitting experimental data

- ***2p calcium recording of Anterolateral Motor Cortex (ALM) with targeted photostimulations on a portion of the trials:*** We train an ensemble of twin RNNs and latent circuits of varying sizes ($n \in [4, 20]$) to deconvolved (with OASIS[4]), bootstrapped neural activity.
- Expectations: Data cannot fully constrain model connectivity; Latent circuits diverge as size increases.

LCs fitted to data align with theoretical results

Smaller latent circuits are comparable at predicting post-perturbation neural responses despite higher validation loss.
*Best 20 models of each size plotted; y-axis limited to

Conclusions

- Low-dimensional data ***do not uniquely constrain*** high-dimensional RNNs, which undermines their interpretability and causal predictive power.
- Trained RNNs connectivity harbors ***unique low-dimensional structure***, which offers maximal predictive power for guiding perturbations.
- When fitting experimental data, validation loss ***does not reflect*** performances in perturbation predictions.

References

1. **Daie, K., Svoboda, K., & Druckmann, S. (2021).** Targeted photostimulation uncovers circuit motifs supporting short-term memory. *Nature neuroscience*, 24(2), 259-265.
2. **Finkelstein, A., Fontolan, L., Economo, M. N., Li, N., Romani, S., & Svoboda, K. (2021).** Attractor dynamics gate cortical information flow during decision-making. *Nature neuroscience*, 24(6), 843-850.
3. **Langdon, C., Engel, T.A. (2025).** Latent circuit inference from heterogeneous neural responses during cognitive tasks. *Nature Neuroscience*, 28, 665–675.
4. **Friedrich, J., Zhou, P., & Paninski, L.** (2017) Fast online deconvolution of calcium imaging data. *PLoS computational biology* 13.3: e1005423